

What happens between signing and the first memo.

A day-by-day breakdown of the kickoff window. For new customers, prospects considering scoping, and any team member who will own the engagement on the customer side.

The kickoff is three weeks of active measurement plus the surrounding setup and delivery. This page covers the full month from signing through the first delivered research memo.

Phase 1: Signing and setup (Days 1-3)

DAY 1

Engagement letter countersigned, invoice issued

Lance sends a one-page engagement letter naming the category, the deliverable shape, the pricing, and the standard terms. You countersign. We invoice the kickoff (\$4,500), due net 15. We do not require payment before measurement begins, but the queue starts moving on signing.

Your side: legal review (typically < 1 day for this length of document) + countersign. Our side: invoice + queue position locked.

DAY 2-3

Query-set scoping call (30 minutes)

One call with you, your designated team contact, and Lance. We bring a starting list of 15-20 queries we think represent how a buyer in your category actually searches. You

bring your knowledge of which queries are commercially most important. We end the call with a frozen, hash-locked query set.

Same call: we lock the competitor cohort. Default for most categories is the 3-5 obvious named competitors plus 1-2 long-tail competitors you watch. Easy to adjust later.

Your side: one person who knows the category · 30 minutes. Our side: Lance · 30 minutes + 1 hour prep.

Phase 2: Measurement window (Days 4-24)

DAY 4

First daily measurement run

Daily measurement begins. Each of the locked queries runs three times per day, on each of the seven AI surfaces we measure: Perplexity, ChatGPT search, Gemini grounded, Microsoft Copilot via Bing, Google AI Overviews (the five citation-grade engines) plus Claude and Gemma (the two model-knowledge engines).

You receive a one-line confirmation email each day for the first week (then weekly thereafter, unless you opt out) saying the run completed with the OK rate. No memo content yet. The first week is establishing baseline.

Your side: nothing required. Our side: scheduled cron + monitoring.

DAY 7

First-week status (optional check-in)

Five days of data captured. If anything looks structurally surprising (a competitor we did not include is dominating the cohort, the query set is producing answers that suggest a wording problem, an engine is rate-limiting), Lance flags it now rather than waiting for the memo.

Most engagements skip this check-in. It exists for the edge cases.

Your side: read a one-paragraph status email if it comes · 2 minutes. Our side: Lance · review the early data.

DAYS 8–21

Continued daily measurement

Same rhythm. Three weeks of data is the minimum cohort size for the pattern-readiness rule to hold (see [/methodology/](#)). Each additional day adds to the signal-vs-noise resolution and the recurring-host classifier output.

Behind the scenes during this window: classification of every cited URL into the nine source-type buckets, competitive cohort tracking against the named competitors, per-engine behavior tracking, anomaly detection if any engine shifts its citation behavior mid-window.

Your side: nothing required. Our side: automated runs + monitoring + analysis prep.

Phase 3: Analysis and memo writing (Days 22-26)

DAYS 22–25

Analysis and draft memo

Lance writes the research memo against the three weeks of captured data. Sections:

- Executive summary (the one thing that matters).
- Methodology and honest scope (what was and was not measured).
- The headline finding (specific to your data, not generic).
- Source-type distribution (which kinds of sources AI cites for your category).
- Per-engine breakdown (where the seven engines differ).
- Competitive gap analysis (which competitors appear on which queries).
- Recurring host analysis (the specific authorities AI engines learned to trust).

- Prepped punch list, ordered by impact, with effort estimates and expected-window-for-movement notes.
- What we did not measure (honest scope limitations).
- Appendices (raw data pointers, classifier rules, query set).

See [/example-engagement/](#) for what this looks like in finished form.

Your side: nothing required. Our side: Lance · ~6-10 hours of focused writing.

DAY 26

Internal grader review

The draft memo runs through the fail-closed factual-and-voice grader before any version reaches you. The grader is a separate AI process that holds the same canonical fact list every customer-facing artifact must honor. If anything in the memo asserts something we cannot substantiate, the grader rejects the draft and Lance revises.

This is the same mechanism that prevents the previous retracted-product claims from re-entering customer-facing work. See the [/retraction/](#) story for why this discipline exists.

Your side: nothing required. Our side: automated grader + Lance revision pass.

Phase 4: Delivery and handoff (Days 27-30)

DAY 27

Memo delivered

You receive the research memo as a PDF (and as markdown, if you prefer that format too) attached to a short cover email naming the headline finding and the highest-priority punch list item. Delivered to your designated recipients only — typically two to four people on your team.

Your raw measurement data (every captured API response, every cited URL with timestamps) is available at your customer data store and is exportable any time.

Your side: read + circulate to your execution team · ~30 minutes for the headline. Our side: Lance · sends.

DAYS 28-30

Optional walkthrough call

If your team wants to discuss the memo synchronously, we schedule a 30-minute walkthrough. Lance walks through the prepped punch list, takes questions, names what we would prioritize and why. If your team is comfortable reading the memo cold, we skip this step.

Your side: optional · 30 minutes. Our side: Lance · 30 minutes.

Phase 5: Ongoing engagement begins (Day 31 onward)

DAY 31+

Continued daily measurement, monthly delta memo

If you opted into the ongoing engagement at scoping, the \$1,500/month retainer takes over. Daily measurement continues on the locked query set. A monthly delta memo arrives on the same calendar day each month, reporting what moved, what did not, and any punch list refreshes based on the new data.

The first delta memo lands roughly 30 days after the kickoff memo. The format is lighter than the kickoff memo: focused on the deltas, with the full punch list status table appended.

You can cancel the ongoing engagement at the end of any month, no claw-back.

Your side: read the monthly memo · ~20 minutes. Our side: automated measurement + Lance · ~3-4 hours per month per category.

What happens if something goes wrong

- **An engine starts rate-limiting or rejecting our queries** mid-window. We surface it in the next status email and adjust the measurement schedule. The deliverable still ships on time, with the honest engine-by-engine OK rate reported.
- **Your query set turns out to need adjustment** after data starts coming in (a query is returning answers about a different category, a query is too broad to be useful). We surface it, propose a revision, and either lock the revised set (with a new hash, restarting the cohort window) or proceed with the original (with the limitation noted in the memo).
- **The cohort-readiness bar is not met** because of API failures (less than 3 usable runs in the window). We extend the window at no additional cost and re-deliver when the bar is met.
- **You disagree with the headline finding.** We walk through the raw data with you. If the data does not support the finding we wrote, we revise the memo. If it does, we may keep the finding and add a section addressing your concern. The data is the data.

What you can expect not to happen

- **No surprise charges.** Kickoff is one fixed price. Retainer is one fixed price.
- **No requests for access to your systems.** If a future engagement requires it, we will tell you up front and you can decline.
- **No upsells during the engagement window.** If a different category would benefit from measurement, we mention it once in the delivery email and never again unless you ask.
- **No PR-style claims of citation lift** in any memo. We report what AI engines cite. We do not claim our work caused the change.

Questions about anything on this page, email lance@neverranked.com.

[Return to NeverRanked](#) · [Example deliverable](#) · [How we measure](#) · [Security posture](#)

