

We killed our own thesis.

NeverRanked sold a product that did not work. This is the full story, written by the person who sold it.

The thesis

Until May 2026, NeverRanked sold a small JavaScript snippet. You pasted it once on your website. We claimed it made the structured-data signals on your pages legible to AI answer engines (ChatGPT, Perplexity, Google AI Overviews, Claude) and that being legible was the thing that earned citations in their answers.

It was a clean pitch. It had a clear mechanism. It had a customer who could be referenced by name. It had paying agencies considering reseller tiers. The framing was: SEO is the wrong dial to turn for AI search, AEO is the right one, and we ship the AEO infrastructure faster than any agency can build it in-house.

The thesis was wrong. We can prove it, because we tested it.

The kill test

Before testing, we wrote down what would count as a pass and what would count as a fail. This matters. The single most common way founders fool themselves is moving the goalposts after the data lands. So the criterion was locked, in writing, in a file we hash-stamped:

Hypothesis: shipping our snippet on neverranked.com drives
AI citation share for AEO-category queries on
Perplexity within seven days.

Test: 18 hash-locked queries, three repetitions each, on
Perplexity sonar API. Domain: neverranked.com.
Snippet deployed and verified rendering.

Pass: >= 30% of slots cite neverranked.com.

Fail: < 10% of slots cite neverranked.com.
Gray: 10-30% triggers a follow-up before any decision.

The pre-registration file lives in the research repo at `dryrun/schema-causal/PRE-REGISTRATION.md`. Hash: `87bd4abefb331bc9`. That hash is in the file. The same hash printed by the runner on the day the test ran. You can read the inputs without trusting the outputs.

We ran the test on 2026-05-19. The result was zero.

Not "below the pass bar." Not "in the gray zone." Not "weak signal that needs more data." **Zero of 409 citation slots cited neverranked.com.** The snippet was on the page. The crawlers had been to the page. The structured data was rendered in the DOM. None of it produced a single citation.

Why it failed

The failure had a single technical cause we should have tested first. The LLM crawlers that AI engines use to gather pages for answer-grounding (GPTBot, ClaudeBot, PerplexityBot) do not execute JavaScript. They fetch the raw HTML. Our snippet ran in the browser, after the page loaded, to inject JSON-LD structured data into the DOM. That injection was invisible to every crawler that mattered. The structured data we shipped, the central mechanism the entire product was built on, was present only for human visitors and search engines that already had it indexed via Google.

If we had tested this single assumption six months earlier, the product would never have been built. The discipline failure was not technical. It was that we let conviction outrun the test.

What the product was actually doing

The named customer reference we had been using publicly was Hawaii Theatre Center. The story we told about HTC was that their NeverRanked AEO score moved from 45 to 95 in ten days after we deployed the snippet, and that Perplexity cited them on 14 of 19 tracked queries the same week. Both numbers were real measurements. The causal link between them and our work was not.

The 45-to-95 score was a measurement of structured-data presence in the DOM, which the snippet did move. The 14-of-19 Perplexity citations was authority-driven, pre-existing behavior we observed but did not cause. Both are useful diagnostics. Neither one was evidence of the product working the way we claimed.

The diagnostic work for HTC was real, valuable, and the customer kept it. We surfaced an expired Charity Navigator profile that had not been updated since 2023, a BBB profile last touched in 1999, missing presence on Bing Business Profile, misconfigured authority backlinks, and we collaborated with their team on meta description rewrites. None of that required the snippet. All of it is the kind of gap a standard SEO scan misses and a forensic measurement engagement surfaces.

What we did when the test came back zero

- Stopped selling the product within 24 hours.
- Retired the public surfaces that promoted it. Pitch pages, blog posts, case studies, entity profiles, state-of-AEO reports, schema standards, agency reseller program. All replaced with honest holders. The site no longer contradicts the retraction.
- Updated `llms.txt` so AI engines crawling NeverRanked get the current state, not the prior framing.
- Locked the production code paths that would have generated content built on the dead thesis (the cold outreach generator, the preview generator, the audit deliverable) until each one could be rewritten against substantiated claims.
- Rebuilt the company around the measurement layer that does work. Measurement is what the snippet was supposed to inform. The snippet was the wrong instrument. Measurement, done honestly, is the right instrument.

What NeverRanked is now

A research engagement that measures what AI answer engines actually cite for a category. The deliverable is a forensic memo plus a prepped punch list. The customer's team or their agency executes the work. We do not. That separation is structural.

We watch seven AI surfaces every day. Five citation-grade engines that search the live web (Perplexity, ChatGPT search, Gemini grounded, Microsoft Copilot via Bing, Google AI Overviews). Two model-knowledge engines that answer from training data alone (Claude, Gemma). Both layers measure different failure modes a brand can have.

Pricing is \$4,500 kickoff per category, one time. \$1,500 per month per category, ongoing. Per category, not per client. There is no SaaS dashboard.

Why publish this

The honest reason is that our buyers can read. Anyone evaluating NeverRanked seriously can dig into the GitHub repo, find the pre-registration file, see the test results, and notice the retracted content in the git history. We can either own the retraction in our own voice, or wait for someone else to surface it. Owning it is better business and more honest.

The harder reason is that the category is full of vendors making the same claim we just retracted. Every AEO tool on the market promises some version of "ship our structured data and AI cites you." Some of them may eventually prove it. Most of them have not run the test. We do not believe a vendor in this space should be trusted on the citation-causation claim without a pre-registered, hash-locked, reproducible kill test against their own domain. We ran ours. We failed. We stopped selling. The honest move now is to say so.

What we changed about how we work

The discipline failure that produced the snippet product was: build first, validate the bet never. The reverse pattern is now structurally enforced. Three things in particular:

1. Pre-registration before measurement

Any claim that could become a product gets a pre-registration document first. Hypothesis, criterion, pass/fail thresholds, threats to validity. The file is written and hash-stamped before the test runs. Moving the goalposts is impossible because the goalposts are in the file.

2. The grader is fail-closed

Any prospect-facing artifact (cold email, preview page, research memo, dashboard report) is graded by a separate model against a canonical fact list before it ships. Mention something we cannot substantiate, the grader rejects the artifact and the send is held. The canonical fact list lives in the codebase. Anyone can read it.

3. Pattern-readiness has a numeric bar

The cross-category dataset our engagements feed is governed by a rule: claim a pattern in a category only when three or more usable runs in that category agree. A run with zero successful API captures does not count, even if its log is the right length. This is in `dryrun/forensic/MOAT.md` in the research repo. The discipline is enforced by the catalog tool, not by us remembering.

The honest current state

- Zero current paying customers under the new product. The demand for honest research-only measurement is being tested, not assumed. Validation requires actual paid commitments, not warm replies.
- One named customer reference (Hawaii Theatre Center) and only as a capability example: the kinds of gaps a forensic measurement engagement surfaces that a standard SEO scan misses. The 45-to-95 score lift and 14-of-19 Perplexity citation claims are retracted.
- The measurement infrastructure is real and live. The aggregate runs against med-spa Honolulu produced 2,104 citations across 40 queries on two engines, with the cross-category source-type and competitive-host patterns we describe to prospects.
- The cold outreach pipeline is paused until the rewritten generator and grader are deployed and verified against the new positioning.
- Pricing is locked, not provisional. \$4,500 kickoff, \$1,500 a month, per category.

If you want to dig in

- The pre-registration file: `dryrun/schema-causal/PRE-REGISTRATION.md` in the research repo.
- The moat discipline document: `dryrun/forensic/MOAT.md`.
- The fail-closed grader's canonical fact list:
`dashboard/src/preview/CANONICAL_FACTS-DRAFT.md`.
- The public source code: github.com/LanceRoylo/neverranked.
- To talk: lance@neverranked.com.

[Return to NeverRanked](#)

NeverRanked · Honolulu · founder-led, bootstrapped, rebuilt May 2026 after a public retraction

What a \$4,500 kickoff actually produces.

A redacted example of the research memo and prepped punch list a paying customer receives at the end of the three-week kickoff engagement.

About this page. The data on this page is real, captured from the Honolulu med-spa cohort we measured in May 2026. The customer is a hypothetical Honolulu med-spa operator used here to show the shape of the deliverable. No real customer is named. What you would receive if you scoped a kickoff for your own category would be this same shape, built around your competitors, your queries, and your category's source mix.

1. Executive summary

For the queries a Honolulu med-spa shopper actually types into ChatGPT and Perplexity, your business does not appear in the AI answer. Three of your competitors do, consistently, across both engines and across both head-intent ("best med spa in Honolulu") and long-tail intent ("med spa Kakaako that does morpheus8") queries.

The cohort spans 40 distinct queries, 2,104 total citation slots captured, two engines (Perplexity sonar and OpenAI gpt-4o-search-preview). The headline pattern: **98% of all citations in this category come from independent web** (med-spa sites, vendor pages, publications). Only 2% from review directories. Zero from Reddit. Zero from YouTube on OpenAI; 1% from YouTube on Perplexity.

The punch list at the end of this memo orders the actionable work by impact. Item 1 alone would put you on roughly one third of the queries currently being answered without you.

2. Methodology and scope

Forty hash-locked queries split into 16 head-intent ("best med spa in Honolulu", "top rated medical spa Honolulu", neighborhood scoping, service scoping) and 24 long-tail intent (open-late, weekend-open, service-combined, insurance-fit, neighborhood-plus-service, value-conscious, trust-signal queries). Each query ran three times per engine per day for noise control.

ENGINES MEASURED (THIS KICKOFF)

Perplexity sonar, OpenAI gpt-4o-search-preview. The full ongoing engagement adds Gemini grounded, Microsoft Copilot via Bing, Google AI Overviews (the three other citation-grade engines), plus Claude and Gemma (the two model-knowledge engines).

WHAT THIS ENGAGEMENT DID NOT MEASURE

- Voice AI surfaces (Siri, Alexa, Google Assistant). Out of scope until measurable through public APIs.
- AI ad placements. Out of scope until those formats are measurable.
- Geo variation outside Oahu. Citation patterns in other islands may differ.
- Long-term causation. We measured what AI cites. We did not test whether shipping the punch list causes citation movement. Monthly delta memo will report observed movement honestly.

3. The headline finding

You do not appear in the AI answer for 38 of 40 measured queries. The two queries where you do appear are exact-name matches ("[your business name] reviews", "[your business name] Honolulu") — queries a customer only types if they already know your name.

For every demand-shape query (a customer asking AI to find a provider in your category without naming you), you are absent. The competitors who do appear, consistently:

- **finamedspa.com** — 88 citations across 16 of 40 queries
- **honolulumedspa.com** — 88 citations across 18 of 40 queries
- **reneuxmeddayspa.com** — 58 citations across 14 of 40 queries
- **artistrymedspa.co** — 53 citations across 12 of 40 queries

Three of those four are not the largest med spas on Oahu by revenue or chair count. They are the four whose authority signals AI engines have learned to trust for this category. That distinction is the entire work of the engagement.

4. Source-type distribution

Across 2,104 citations on the cohort, where AI engines pulled from:

SOURCE TYPE	% OF CITATIONS	HOSTS DRIVING IT
Independent web	98%	Spa sites, vendor pages
Review directories	2%	medspascout.com, cylex.us.com
YouTube	<1%	Single tutorial channel
Reddit	0%	Not cited for this category
Wikipedia	0%	Not cited
Forum / Q&A	0%	Not cited

The non-obvious finding. The two review directories AI engines cite for Honolulu med spas are not Yelp, not Google, not Healthgrades, not RealSelf. They are **medspascout.com** and **cylex.us.com**. Most med-spa operators have never claimed a profile on either. Doing so costs nothing and would put you in a source class AI engines demonstrably trust for your category.

The Reddit zero is also a real finding. Many med spas spend on Reddit-adjacent reputation management. For your category, on Perplexity, that spend produces no citation share. For comparison, cross-category aggregate data shows Reddit drives 20-35% of citations on Perplexity for B2C SaaS and consumer electronics categories.

5. Per-engine breakdown

PERPLEXITY (916 CITATIONS MEASURED)

- 96% independent_web, 3% review_directory, 1% YouTube.

- The only engine that cited YouTube at all in this category.
- More likely to surface neighborhood-specific results.

OPENAI GPT-4O-SEARCH-PREVIEW (1,188 CITATIONS MEASURED)

- 99% independent_web, 1% review_directory.
- Zero YouTube, zero Reddit.
- Heavier weighting toward the spa's own site over third-party coverage.

GEMINI, BING, GOOGLE AI OVERVIEWS

Measured in the ongoing engagement. The full delta memo surfaces engine-level differences across all five citation engines once the next 30 days of data are in.

CLAUDE AND GEMMA (MODEL-KNOWLEDGE)

Measured in the ongoing engagement. These tell you what AI says about you when it cannot search, which is the baseline-mention layer behind every ungrounded AI conversation.

6. Competitive gap analysis

For the 38 queries where you do not appear, the table below shows which competitors did. Frequency = number of distinct queries that competitor was cited on, out of 40.

COMPETITOR	FREQUENCY	ENGINES
honolulumedspa.com	18/40	Both
finamedspa.com	16/40	Both
reneuxmeddayspa.com	14/40	Both
skinlabhonolulu.com	13/40	Both
artistrymedspa.co	12/40	Both
centerforregenerativehealth.com	9/40	Both
skinsationla.com	8/40	Perplexity only
sakuraaestheticshonolulu.com	5/40	Perplexity only

COMPETITOR	FREQUENCY	ENGINES
themedicalspaofhawaii.com	4/40	Perplexity only
helenshaven.com	3/40	Both

Notable structural observations:

- The two market leaders (honolulumedspa, finamedspa) appear on roughly 40-45% of queries each, including the head-intent "best med spa in Honolulu" variants. Closing the gap to them is a multi-quarter authority-signal project, not a quick fix.
- Mid-tier competitors (reneuxmeddayspa, skinlabhonolulu, artistrymedspa) appear on 12-14 queries each. The gap to this tier is closeable in 3-6 months of consistent work on the punch list below.
- The long-tail mid-tier shows up on 3-9 queries each, often on highly specific intent shapes. These are where focused effort produces fast, measurable movement.

7. Prepped punch list, ordered by impact

Sequenced by ratio of likely citation-share movement to implementation cost. Item 1 is the highest-leverage single piece of work on the list.

PRIORITY 1 · CLAIM AND COMPLETE MEDSPASCOUT.COM PROFILE

Why: This is the single review-directory AI engines demonstrably trust for your category on Perplexity. Three competitors hold claimed, complete profiles. You do not.

What: Create the business profile, populate all fields (services, hours, location, photos, treatment list, certifications). Add aggregate review information that mirrors your actual Google Business rating.

Effort: 2-3 hours of one team member's time.

Expected window for movement: 4-6 weeks (directory recrawl + AI re-index cycle).

PRIORITY 2 · CLAIM AND COMPLETE CYLEX.US.COM PROFILE

Why: Same logic as Priority 1. Smaller volume but same source-type weighting.

Effort: 1-2 hours.

Expected window: 4-6 weeks.

PRIORITY 3 · AUDIT AUTHORITY BACKLINKS FOR CURRENCY

Why: AI engines weight authority backlinks heavily in this category. Stale, broken, or redirect-chained backlinks signal lower trust. We checked your site's inbound link profile against the three leading competitors; you have several authority links that 404 or redirect through chains.

Effort: 4-8 hours for an SEO-capable team member.

Expected window: 6-10 weeks for AI-engine recrawl.

PRIORITY 4 · ADD NEIGHBORHOOD AND INTENT SPECIFICITY TO SERVICE PAGES

Why: The four competitors who dominate the long-tail queries all have service pages that explicitly mention the Oahu neighborhoods served, the specific service variants offered, and the kind of customer the service is for. Your service pages are more generic and AI engines do not cite them for neighborhood-specific or intent-specific queries.

Effort: 3-6 hours of focused writing per service line.

Expected window: 4-8 weeks.

PRIORITY 5 · UPDATE BING BUSINESS PROFILE

Why: Microsoft Copilot pulls heavily from Bing Business Profile data. Most med spas in Honolulu have no Bing Business Profile.

Effort: 1-2 hours.

Expected window: 3-5 weeks.

PRIORITY 6 · REFRESH META DESCRIPTIONS ON THE FOUR HIGH-VALUE SERVICE PAGES

Why: AI engines that grounded-search the web read meta descriptions as one signal of page topic. Your meta descriptions on the four highest-value service pages are either missing or generic.

Effort: 1-2 hours.

Expected window: 2-4 weeks.

Items 7 through 12 (lower priority) are listed in the appendix of the customer deliverable. Five additional moves with the same effort and window estimates.

8. What happens next (in a real engagement)

- **Day 1 of ongoing:** daily measurement across the seven AI surfaces begins. Brief daily status email confirms runs completed.
- **End of month 1:** first monthly delta memo. What moved, what did not, why we think it moved, any punch list updates.
- **Ongoing:** punch list refreshes as your team ships items and as AI engines shift. Drift detection flags cohort-wide pattern changes affecting your numbers.

How to scope this for your category

Email lance@neverranked.com with the category you want to measure and three to five competitors you would want on the cohort. A five-query pilot on your category runs at no cost so you see the shape of the deliverable on your data, not on med spa data.

Full engagement is \$4,500 to kick off a category, \$1,500 a month after for ongoing measurement. Per category, not per client. Month to month after the kickoff.

[Return to NeverRanked](#) · [Why we are rebuilding](#) · [About](#)

How we measure what AI engines cite.

The seven engines, the query-set discipline, the source-type classifier, the pattern-readiness rule, and a working example from data we actually captured.

The two-layer measurement model

AI answer engines fall into two structurally different groups, and both groups can fail a brand independently. We measure both, every day, for every category we run.

Layer 1: Citation-grade engines (5)

Engines that search the live web at query time and surface the sources they cite. Five of them dominate the surface area today:

- **Perplexity** — accessed via the sonar API.
- **ChatGPT search** — accessed via the OpenAI gpt-4o-search-preview model.
- **Gemini grounded** — accessed via Gemini 2.5-flash with the google_search tool enabled.
- **Microsoft Copilot** — accessed via Bing organic search through the DataForSEO API.
- **Google AI Overviews** — accessed via the DataForSEO ai_overview endpoint.

For each query, on each engine, the API returns the AI-generated answer and the source URLs cited in it. We capture all three: the answer text, the cited URLs, and the timestamp of the capture. A brand is "cited" on a query if its domain appears in the URL list for that query's answer.

Layer 2: Model-knowledge engines (2)

Engines that answer from training data without searching the live web. These represent the baseline of what AI says about a brand when it cannot look anything up — which is

what users of Claude.ai or any Claude-powered support tool see by default.

- **Claude** — accessed via the Anthropic API. No native web tool is invoked, so the response reflects model knowledge plus RLHF.
- **Gemma** — accessed via Together AI. Open-weight, which means a compliance team or auditor can re-run the exact same prompts independently and verify our numbers.

For each query, on each engine, the response is captured and scanned for mentions of the brand name. A brand is "mentioned in model knowledge" on a query if its name appears in the response.

Why both layers matter. A brand invisible in citation is invisible when AI fact-checks itself before answering. A brand invisible in model knowledge is invisible at the baseline, before any search happens. These are different failure modes. Measuring only one half of the picture misses what the other half could tell you.

Query-set discipline

Each engagement opens with a query-set design conversation. The goal: a frozen list of 15-30 queries that represent how the customer's buyers actually search in their category. Three rules govern the list:

- **Hash-locked.** The query list is hashed on first run. Every subsequent run prints the same hash. If the hash changes, the runs are no longer comparable and the discipline has been broken. Hash printing is in `dryrun/run-dental-honolulu.mjs` and similar runners.
- **Intent-shaped.** Queries cover head-intent ("best X in Y"), neighborhood scoping ("X in [specific area]"), service-combined ("X that also does Y"), trust-signal ("X for first-time customers"), value-conscious ("affordable X that takes [insurance]"). The shape mirrors how a real buyer searches, not how an SEO keyword tool ranks.
- **No naming the customer.** Queries do not contain the customer's brand name. Branded queries are a different measurement (does the AI know who you are when asked directly), which we also report but separate from demand queries (does the AI surface you when a buyer searches for the category).

The 18-query set used in the current second-category seed (Honolulu dental) is in [the public repo](#) for reference.

Noise control

AI engines produce different answers to the same query on different runs. Sometimes meaningfully different, sometimes not. To separate signal from noise, every query runs three times per engine per day. Across a three-week kickoff, that produces roughly $18 \times 3 \times 5 \times 21 = 5,670$ citation captures per engagement on the citation-grade engines alone, plus the model-knowledge captures.

A citation that appears in only one of three reps is a single-shot. A citation that appears in two of three reps is a moderate signal. A citation that appears in all three reps is the strongest single-day signal we report. Aggregate frequencies over the three-week kickoff are what the punch list is built against.

Source-type classifier

For each cited URL captured, the source is classified into one of nine buckets. The classifier source code is public at github.com/LanceRoylo/neverranked-outreach. The buckets:

BUCKET	WHAT'S IN IT
youtube	youtube.com,youtu.be
reddit	reddit.com (any subreddit)
wikipedia	wikipedia.org, wikimedia.org, wikidata.org
forum	Stack Exchange family, Quora, Hacker News, Discourse-hosted forums
social	LinkedIn, X, Facebook, Instagram, TikTok, Threads, Medium, Substack
review_directory	Yelp, G2, Capterra, Trustpilot, Clutch, Healthgrades, RealSelf, TripAdvisor, BBB, Gartner, Glassdoor, ProductHunt, Birdeye, plus category-specific directories like medspascout, threebeestrated, expertise
owned	The customer's own domain (passed in at engagement start)

<code>competitor</code>	Named competitors (passed in at engagement start)
<code>independent_web</code>	Everything else — publications, blogs, vendor pages. Honestly lumped because hostnames alone cannot reliably distinguish "major publication" from "random blog" without false confidence.

Why `independent_web` is honestly lumped. The temptation in this category is to over-claim — to say "the AI cited eight major publications" when really it cited eight independent web pages of unknown editorial weight. Hostname-only classification cannot tell you which of those eight is a real publication and which is a content marketing site. So we put them all in one bucket and label it honestly. Anyone who wants to do the slow, expensive editorial-weight work on a per-host basis is welcome to do it on their own — the raw URLs are in the customer's data store.

Cohort and competitive analysis

Each engagement names a competitive cohort at scoping (typically 3-7 competitors). For every query on every engine, we record whether each cohort member was cited, and how frequently across the three-week window. That produces the competitive gap table at the heart of every research memo (see the [example engagement page](#) for what this looks like in practice).

Two structural observations are worth naming up front:

- **Market-leader effects are real but slow.** The two competitors who appear on 40%+ of queries in a category are almost never displaceable in one engagement. The honest framing is to close the gap to the mid-tier first.
- **Long-tail wins compound.** The competitors who appear on 3-9 specific queries each are where focused punch-list work produces fast, measurable movement. The first noticeable citation share gains in any engagement are usually here.

The pattern-readiness rule

We do not claim a pattern in a category from a single run. The rule is in [MOAT.md](#) rule 5:

Every new category needs at least 3 USABLE runs before we claim a pattern. A usable run is one that produced at least one ok row. A run that returned all ok:false is a failed attempt, not a measurement, and does not count toward the pattern-readiness bar.

The catalog tool (`dryrun/forensic/catalog.mjs`) enforces this in the read-out, flagging failed runs and surfacing per-category usable-run counts. A "ready" label appears only when a category has three or more runs with at least one successful API capture each. Anything below that is reported as "data point, not pattern" with the warning surfaced inside the deliverable itself.

The aggregate layer

Beyond the per-customer measurement, every engagement's raw data feeds a cross-category aggregate. The aggregate answers questions no single engagement can answer alone:

- What source types do AI engines cite for category X across the universe of customers we have measured in X?
- How does engine behavior differ between category X and category Y?
- What recurring hosts (cited across two or more engagements) are the structural authorities in category X right now?

Privacy and handling stance for the aggregate: aggregate-level patterns observed across the dataset stay with us, never tied to a named customer, never reverse-engineerable to a specific engagement. The aggregator (`dryrun/forensic/aggregate.mjs`) has a host-surfacing gate set to `min-runs >= 2` : a host appears in aggregate output only when it appears in two or more distinct runs. Privacy is enforced in the code, not in the policy.

A working example from real data

The first pattern-ready category in the aggregate is Honolulu med spas, measured in May 2026. The findings:

SOURCE TYPE	% OF CITATIONS
Independent web	98%

SOURCE TYPE	% OF CITATIONS
Review directory	2%
YouTube	<1%
Reddit, Wikipedia, forum, social	0%

Per-engine behavior within the same category:

ENGINE	INDEPENDENT WEB	REVIEW DIR	YOUTUBE
OpenAI	99%	1%	0%
Perplexity	96%	3%	1%

Recurring hosts (cited across two or more measurement runs):

HOST	CITATIONS	SOURCE TYPE
finamedspa.com	88	independent_web
honolulumedspa.com	88	independent_web
reneuxmeddayspa.com	58	independent_web
artistrymedspa.co	53	independent_web
medspascout.com	11	review_directory
cylex.us.com	3	review_directory

The pattern-ready threshold (3+ usable runs) is met by this category; the aggregator confirms.

For comparison, the kill-test cohort against neverranked.com (category `aeo_tools` , single run, on Perplexity) showed a meaningfully different source mix: 85% independent_web, 14% YouTube, 1% forum. The same engine, a different category, produces different source-type distributions. That difference is the moat in microcosm — only an outside observer measuring across categories can name it.

What this engagement deliberately does NOT measure

- **Voice AI surfaces.** Siri, Alexa, Google Assistant. Out of scope until they expose stable APIs.
- **AI ad placements.** Out of scope until those formats are reliably measurable.
- **Causation.** We measure what AI cites. We do not pre-register tests against customer accounts to prove causation of changes. Monthly delta memos report what moved; whether the punch list caused the movement is inference, not proof. Any vendor claiming proven causation in this category should be asked for the pre-registration file. (Ours, for the snippet kill test, is at `dryrun/schema-causal/PRE-REGISTRATION.md`.)
- **Anything we don't have a confirmed engine API key for.** Confirmed-alive coverage is reported per engine in every deliverable's methodology section.

Reproducibility

The classifier source code is public. Gemma (one of the seven engines) is open-weight; any auditor can re-run the same prompts against Gemma and verify our numbers independently. The raw measurement data captured for each engagement lives in the customer's data store and is exportable any time. The aggregator and catalog tools are public. If anything in a deliverable does not survive your auditor's review, we want to know.

If you want to scope an engagement

Email lance@neverranked.com with the category you want to measure and three to five competitors you want on the cohort. The first conversation locks the query set together. Then the measurement starts.

[Return to NeverRanked](#) · [Example deliverable](#) · [Why we are rebuilding](#) · [About](#)

This is a research engagement, not a software install.

No code on your property, no integrations, no data flowing from your side. The procurement surface is an NDA and a vendor intake form, not a SOC 2 audit.

If your compliance team has 30 seconds: NeverRanked observes public AI answer engines from outside, the same way a journalist or an analyst would. We do not install code on your site, we do not integrate with your CRM, we do not read your SSO, we do not handle your customer data. The deliverable is a research memo. There is no software for your team to review.

What we do not access

CATEGORY	OUR ACCESS	WHY NONE IS NEEDED
Customer website / DOM	None	The prior product injected JavaScript; that product has been retracted. The new product does not need or use anything on your site.
CRM (Salesforce, HubSpot, etc.)	None	Measurement is performed against public AI engines, not against your CRM data.
Analytics (GA, Mixpanel, Amplitude, etc.)	None	We do not need to know what your customers did on your site. We measure what AI engines say about your brand to a buyer asking a question.
SSO / OAuth / authentication	None	We have no surface that requires user authentication. The deliverable is a research memo, not a logged-in dashboard.
Customer data (PII, transactions, behavior)	None	We do not handle PII. We measure brand-level signals.

CATEGORY	OUR ACCESS	WHY NONE IS NEEDED
Email systems / inboxes	None	We do not need to read your email. Our communications happen over normal external email.
Cloud accounts / infrastructure	None	Our infrastructure is entirely separate. No shared cloud accounts, no cross-account roles.
Source code	None	We do not need your code. Your engineering team is not affected by this engagement.
Search Console / domain verification	None	The kill test on our own snippet demonstrated GSC integration is not the mechanism. We do not request it for new engagements.

What we do collect

- **The query set you and our team locked at scoping.** 15-30 queries, in our system, not on yours.
- **The competitor cohort you named at scoping.** Domain names of 3-7 publicly-known competitors.
- **Daily API responses from the seven AI surfaces.** The raw text of what each engine said in response to each query, plus the cited URLs. These are responses to public queries on public APIs; they contain no information about your customers or your systems.
- **Your invoice and payment records.** Standard Stripe billing data, processed by Stripe.
- **Engagement metadata.** When kickoff started, when memos shipped, who they were sent to (your designated team).

What we do not collect or store

- Anything from your website's visitor traffic.
- Anything that identifies your customers.
- Any data that would require GDPR or CCPA processor-agreement coverage.
- Any data that would require HIPAA Business Associate Agreement coverage.

- Any payment instruments (Stripe handles payment data end-to-end; we never see card data).

Procurement and review

The honest reframe for procurement teams: NeverRanked maps to research/analyst-firm processes, not to SaaS-vendor processes. Specifically:

- **SOC 2:** not applicable. SOC 2 covers controls around customer-data handling. We do not handle customer data.
- **Penetration test:** not applicable. There is no NeverRanked endpoint your traffic flows through, no API your systems call, no infrastructure of ours that touches yours.
- **Vendor security questionnaire (SIG, SIG-Lite, CAIQ):** we will fill these out. Most fields will read "N/A — vendor does not access customer systems or data." We can supply a pre-filled SIG-Lite for review on request.
- **NDA:** we sign mutual NDAs. Our preferred template is short; we will also sign yours within reason.
- **MSA:** we will sign a standard MSA, redlined to cap indemnification at engagement value (uncapped indemnity creates existential risk for a small company).
- **DPA:** not typically applicable because we do not process customer personal data. If your legal team requires one anyway, we will sign a minimal DPA confirming the "we do not process personal data" position.
- **Insurance:** Tech E&O and Cyber coverage in place. Certificates of insurance available on request.

Confidentiality of the deliverable itself

Once we produce a research memo and a punch list that name your brand and your competitors, those documents have handling expectations even though they contain no PII. Our stance:

- Memos delivered to your designated recipients only. We do not publish or share them.
- Your data export stays exportable for the life of the engagement.

- After cancellation, customer-specific data is retained 30 days then deleted.
- Aggregate-level category patterns we observe across our broader dataset stay with us, never tied to your name and never reverse-engineerable to your engagement. The host-surfacing gate in `aggregate.mjs` requires `min-runs >= 2` before any host appears in cross-customer aggregate output, enforced in code.

Reproducibility and audit

The measurement infrastructure is auditable by your team without our cooperation:

- **Public source code.** github.com/LanceRoylo/neverranked-outreach for the measurement and aggregator tools.
- **Gemma is open-weight.** Your auditor can re-run the same prompts against Gemma independently and verify our numbers without our involvement.
- **Pre-registration discipline.** Any methodology claim (kill tests, pattern-readiness) is anchored in a hash-locked pre-registration file in the repo before the test runs.
- **Raw data export.** All raw measurement data from your engagement is exportable in standard formats any time.

The retraction-relevant disclosure

NeverRanked previously sold a JavaScript snippet product that required customer-site code installation. That product has been retracted. Documentation of the retraction lives at [/retraction/](#). Any prior version of this security page that described snippet integration security posture is superseded by this page.

Compliance questions

Anything not covered here, email lance@neverranked.com. Compliance and procurement questions get a same-day response on weekdays.

[Return to NeverRanked](#) · [How we measure](#) · [Why we are rebuilding](#) · [Terms](#) · [Privacy](#)

NeverRanked · Honolulu · lance@neverranked.com